

# accusonus






## **Audio Source Separation**

Elias Kokkinis, CTO  
[elias@accusonus.com](mailto:elias@accusonus.com)

Accusonus is a Greek start-up, focusing on **innovative digital audio technologies**. The company's mission is to offer advanced solutions and unique business services for the music and speech technology sectors.

## Board of Directors

	<b>Alexandros Tsilfidis, PhD</b> Founder & CEO	<ul style="list-style-type: none"> <li>•Electrical and Computer Engineer</li> <li>•MPhil and PhD in digital audio</li> </ul>
	<b>Elias Kokkinis, PhD</b> Founder & CTO	<ul style="list-style-type: none"> <li>•Electrical and Computer Engineer</li> <li>•PhD in digital audio</li> </ul>
	<b>Michael Tzannes, PhD</b> Board member and advisor	<ul style="list-style-type: none"> <li>•Founder and ex-CEO of Aware Inc (Nasdaq: AWRE)</li> <li>•Founder and Executive Manager at Tzannes Patent management LLC</li> </ul>

**+ a great team!**



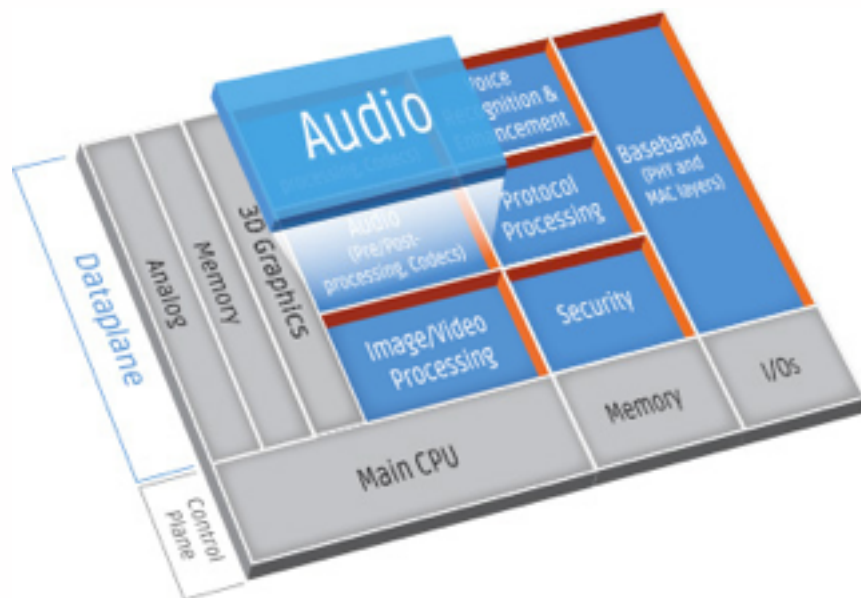


## Focus-BDR/Focus-MDR

- Patent-pending
- The first dereverberation product in the market
- The first dual channel dereverberation in the market
- Combines state-of-the-art dereverberation algorithm with room acoustics

## Focus-DNR

- State-of-the art denoise algorithm reengineered for tight integration with dereverberation



cā dence™



[https://github.com/EliasKokkinis/  
audio-source-separation](https://github.com/EliasKokkinis/audio-source-separation)



# **Audio Source Separation**

## **Part I - Why do we need it?**

- ▶ Can you understand what is being said?



- ▶ Can you focus on the piano?
- ▶ Can you focus on the guitar?





## Man vs. Machine

- ▶ Would a computer be able to do it?
- ▶ Probably **no!**
- ▶ We need audio source separation to *help computers understand* auditory events
- ▶ Applications:
  - speech enhancement and ASR performance improvement
  - music information retrieval applications (MIR)
  - better audio in studio and live (**that's what we do!**)
  - and many more...

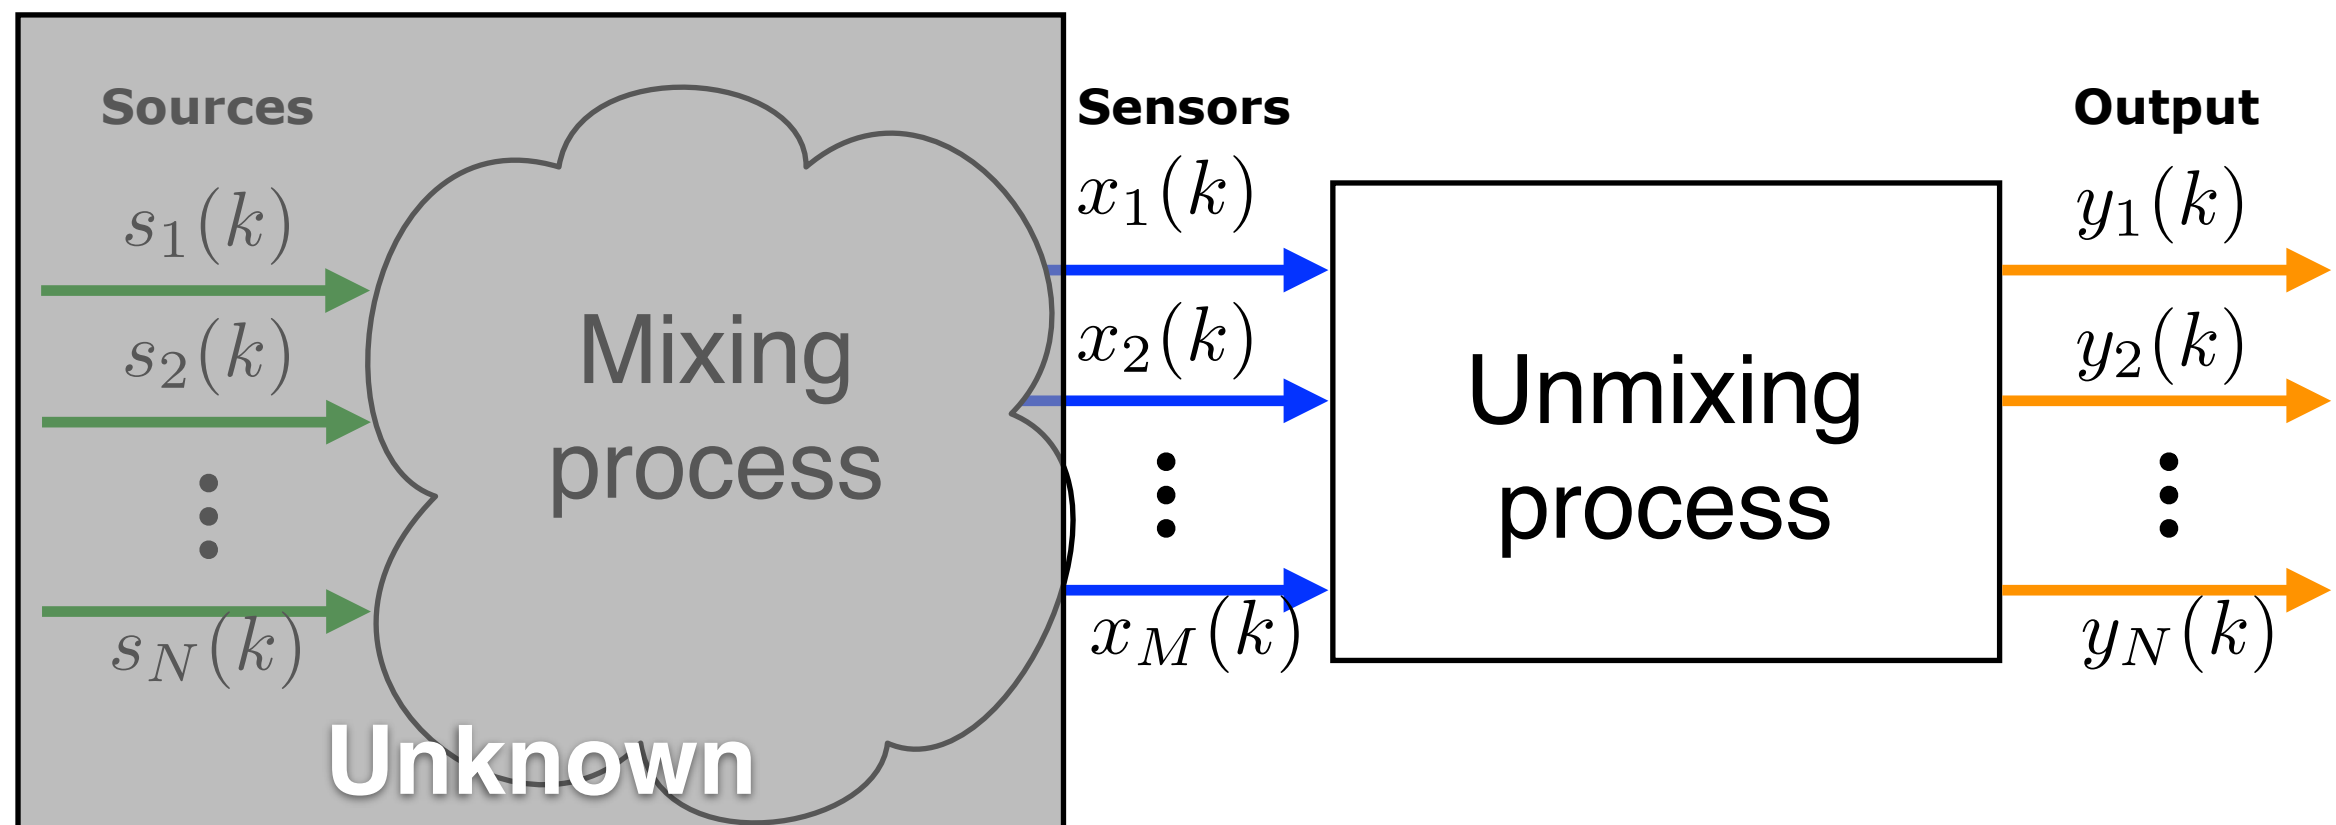
# **Audio Source Separation**

## **Part II - Blind source separation**



## Problem formulation

- Consider **N** sources and **M** sensors (microphones)



## Problem formulation

- ▶ Consider **N** sources and **M** sensors (microphones)
- ▶ For each time instant

$$\mathbf{x}(t) = \mathbf{A}\mathbf{s}(t)$$

- ▶ **A** is the MxN **mixing matrix**
  - instantaneous mixing
  - three distinct cases:
    - overdetermined ( $M > N$ )
    - determined ( $M = N$ )
    - underdetermined ( $M < N$ ). when  $M = 1$  things get really tough...



## Problem formulation

- ▶ Consider **N** sources and **M** sensors (microphones)
- ▶ For each time instant

$$\mathbf{x}(t) = \mathbf{A}\mathbf{s}(t)$$


- ▶ **Blind** source separation
  - estimate an **unmixing matrix** **W** such that

$$\hat{\mathbf{s}}(t) = \mathbf{W}\mathbf{x}(t)$$

- with **only**  $\mathbf{x}(t)$
- we **don't** know anything about **A** or **s(t)**

## Problem formulation

- ▶ The problem **cannot** be solved completely blind
- ▶ We have to make some **assumptions** about the source signals
- ▶ BSS was first introduced in the 1980s
- ▶ A very active research field for more than 20 years





- ▶ Main assumption: **statistical independence (SI)**
  - the most widely used assumption
  - leads to Independent Component Analysis (ICA)
- ▶ Intuition:
  - knowing something about one signal does not give you any information about another signal
  - this is violated when signals are mixed

- ▶ ICA methods estimate an unmixing matrix so that the estimated signals are as SI as possible
- ▶ *How can we measure statistical independence?*
  - non-Gaussianity
  - mixed SI signals tend towards Gaussian (via CLT)
- ▶ *How can we measure non-Gaussianity?*
  - kurtosis: how “spiky” is a pdf
    - 4th order statistic, sensitive to outliers
  - negentropy: difference between the entropy of Gaussian and a given distribution
    - difficult to calculate



- ▶ We want to maximize a measure of non-Gaussianity

$$\mathcal{J}_G = \sum_{i=1}^N \mathbb{E}\{G(\mathbf{w}_i^T \mathbf{x})\}$$

- ▶ G is a non-linear function
  - this function defines the measure that is maximized
- ▶ FastICA: an iterative method to maximize this cost function

Hyvärinen, A.; Oja, E. (2000). ["Independent component analysis: Algorithms and applications"](#).  
*Neural Networks* **13** (4–5): 411–430

- ▶ The update for FastICA is

$$\mathbf{w}_i = E\{\mathbf{x}g(\mathbf{w}_i^T \mathbf{x})\} - E\{g'(\mathbf{w}_i^T \mathbf{x})\}\mathbf{w}_i$$

- $g()$  is the derivative of  $G()$
- $g'()$  is the derivative of  $g()$

- ▶ For kurtosis maximization

$$g(y) = y^3 \qquad g'(y) = 3y^2$$

- ▶ For negentropy maximization

$$g(y) = \tanh(y) \qquad g'(y) = 1 - (\tanh(y))^2$$

► Implementation issues

- Orthonormalization of  $\mathbf{W}$  after each iteration

$$\bar{\mathbf{W}} = \mathbf{W} \left( \mathbf{W}^T \mathbf{W} \right)^{-\frac{1}{2}}$$

- Preprocessing

- Center the data (remove the mean)
- Whiten the data: uncorrelated data with unity variance (use eigenvalues)



- ▶ Ideally

$$\mathbf{W}\mathbf{A} = \mathbf{I}$$

- ▶ BSS suffers from a set of indeterminacies
- ▶ Formally this is stated as

$$\mathbf{W}\mathbf{A} = \mathbf{P}\mathbf{\Lambda}\mathbf{I}$$

- permutation matrix
- scaling matrix (diagonal)
- identity matrix

$$\mathbf{\Lambda} = \mathbf{P} \begin{bmatrix} \lambda_1 & 0 & 0 & 0 \\ 0 & 1 & \lambda_2 & 0 \\ 0 & 0 & 0 & 1 \\ 0 & 0 & 1 & \lambda_3 \end{bmatrix}$$

- ▶ Permuted and scaled signals are still SI!




## Convolutional mixing

- ▶ Instantaneous mixing does not actually occur!
  - but it is very useful of biomedical applications
- ▶ Audio sources are typically inside rooms!
  - The mixing matrix becomes a set of filters
  - The number of parameters to estimate increases drastically!
- ▶ Convolutional BSS: a much harder problem!



# Convolutional mixing

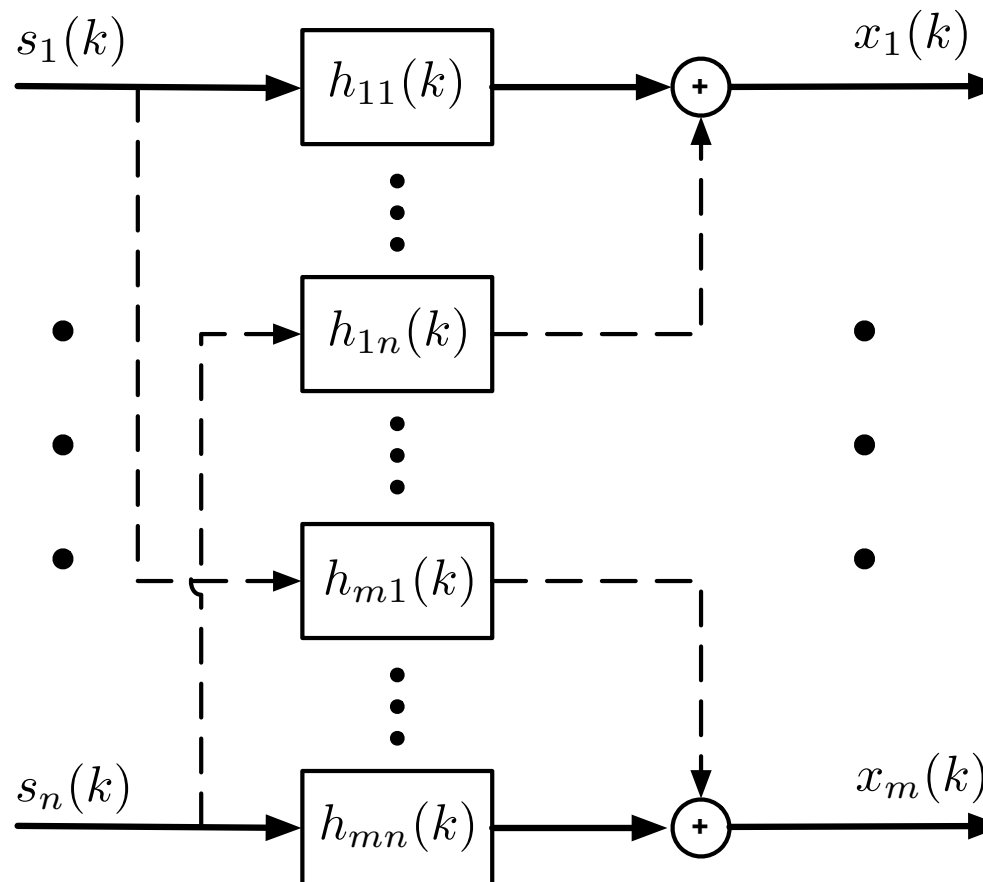
- The room as a filter



- ▶ The room as a filter
  - Modeled as an FIR filter
  - The effect on a signal: convolution
  - The number of coefficients is related to  $f_s$  and RT60
  - The room impulse response (RIR) characterizes a room
    - for the specific source-microphone position!
  - RIRs are non-minimum phase
    - their inversion is not straightforward!

► Convolutional mixing

- A block diagram



- or in equation form

$$x_i(t) = \sum_j \sum_k a_{ij}(t) s_j(t - k)$$



## Freq. domain ICA

- ▶ An elegant solution: transform to frequency domain
  - convolution becomes a multiplication

$$\mathbf{X}(\omega) = \mathbf{A}(\omega)\mathbf{S}(\omega)$$

- ▶ Solve an instantaneous ICA problem for each bin



- ▶ But: ambiguities need to be resolved!
  - the **permutation** problem is quite complex but can be solved
  - **scaling** ambiguities introduce spectral distortions
- ▶ Further issues:
  - the length of the FFT has to be greater than the RIR length for the multiplication assumption to hold (due to **circular convolution**)
  - however a long FFT increases the number of parameters to adjust and **slows convergence**
  - also for longer windows the SI assumption **collapses**

- ▶ Convolutional BSS methods are quite involved
  - beyond the scope of this workshop
- ▶ They suffer from significant limitations in realistic audio applications



# References

- ▶ Two very good books



- ▶ Review chapter “*Convolutive Blind Source Separation Methods*” by M. Pedersen et al. in “Handbook of Speech Processing” (Springer)

# **Audio Source Separation**

## **Part III - Non-negative matrix factorization**



- ▶ Non-negative matrix factorization (NMF)
  - Given a non-negative matrix  $\mathbf{V}$ ,
  - find two matrices  $\mathbf{W}$  and  $\mathbf{H}$  such that

$$\mathbf{V} \approx \mathbf{W}\mathbf{H}$$


$$\mathbb{R}_+ : \{x \in \mathbb{R}, x \geq 0\}$$

$$\mathbf{V} \in \mathbb{R}_+^{F \times N}, \mathbf{W} \in \mathbb{R}_+^{F \times K}, \mathbf{H} \in \mathbb{R}_+^{K \times N}$$



## A bit of history

- ▶ First introduced by Paatero and Tapper (1994)
- ▶ Became widely known with Lee and Seung (1999)
- ▶ Google Scholar reports more than 132.000 results
- ▶ A very strong research field even after 15 years.



## The matrices (1)

- ▶ **V** of size  $F \times N$  is a collection of column vectors  $\mathbf{v}_n$ 
  - each vector is a **data point**
  - each vector is  $F$ -dimensional (or it has  $F$  features)
  - there are  $N$  data points in total

$$\mathbf{V} = [\mathbf{v}_1 \ \mathbf{v}_2 \ \cdots \ \mathbf{v}_N]$$

$$\mathbf{v}_n = \begin{bmatrix} v_{1n} \\ v_{2n} \\ \vdots \\ v_{Fn} \end{bmatrix}$$

## The matrices (2)

- ▶ **W** is the basis matrix of size  $F \times K$
- ▶ **W** has  $K$  columns of size  $F$
- ▶ The columns of **W** span a linear space



## The matrices (3)

- ▶ **H** is the weight matrix of  $K \times N$
- ▶ **H** has  $N$  columns of size  $K$ 
  - The  $n$ -th column of **H** represents the coordinates of the  $n$ -th data point in the space defined by **W**

$$\mathbf{v}_n \approx \mathbf{W} \mathbf{h}_n$$

- ▶ **K** is the **rank** of the factorization
  - Typically  $K$  is chosen such that  $K \ll FN$
  - This leads to a dimensionality reduction in the data
  - For  $K > F$ , matrix **W** becomes overcomplete.
  - Overcomplete basis matrices are related to sparse signal processing.
  - $K$  is the number of “building blocks”

# Terminology & Notation

- ▶ A note on terminology:
  - $\mathbf{W}$  is called basis matrix
    - The columns of  $\mathbf{W}$  are called basis or *atoms*
    - $\mathbf{W}$  is also called a dictionary (in sparse signal processing)
  - $\mathbf{H}$  is called the weight or *gain* matrix
  - $K$  is the *rank* of the factorization (also called number of components)



## NMF discovers parts!

- ▶ **Non-negativity** leads to **parts-based** decomposition
  - The “building blocks” of  $\mathbf{W}$  must be combined **constructively**
  - The algorithm is forced to discover the parts that make up  $\mathbf{V}$

- ▶ Some applications of NMF
  - image processing
  - text mining
- ▶ We are concerned with audio applications.
- ▶ *How can we interpret  $V$ ,  $W$  and  $H$  in audio?*

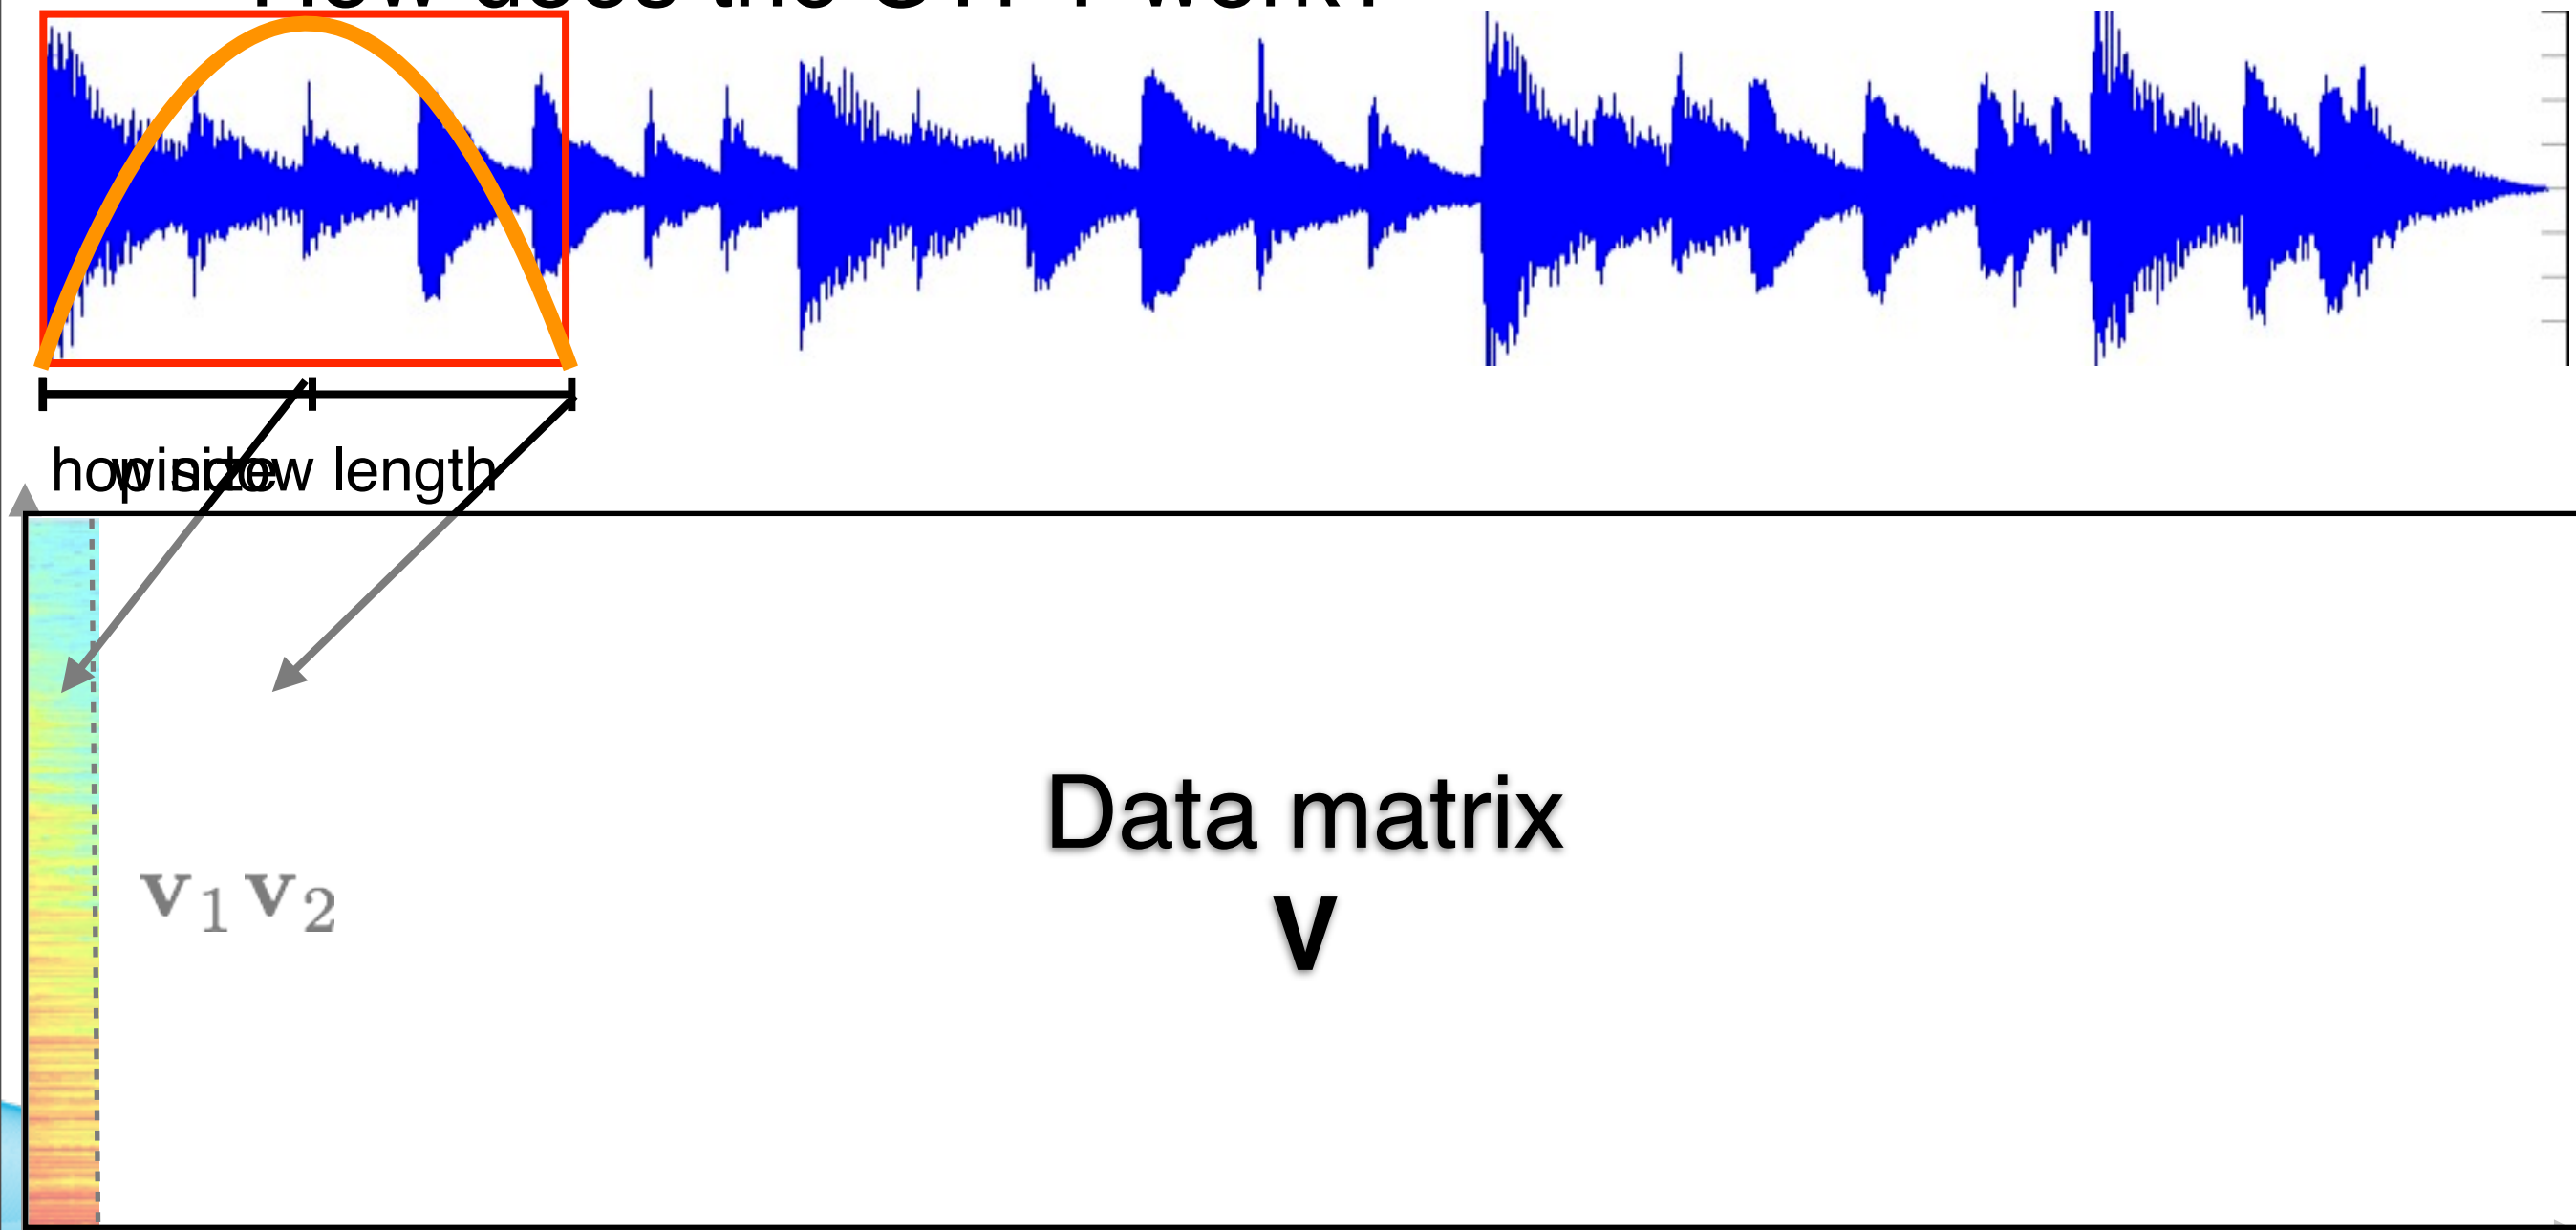
## NMF in audio context

- ▶ **V** represents the audio spectrogram
  - F is the number of frequency bins
  - N is the number of audio frames
- ▶ **W** consists of *spectral profiles*
  - i.e. the magnitude spectra of elementary sources in **V**
- ▶ **H** represents the *gains or activation functions* of the spectral profiles
  - i.e. in a given frame which elementary sources are active



# The spectrogram

- ▶ The spectrogram is the magnitude of the STFT
- ▶ How does the STFT work?




## NMF in audio context

- ▶ **V** represents the audio spectrogram
  - F is the number of frequency bins
  - N is the number of audio frames
- ▶ **W** consists of *spectral profiles*
  - i.e. the magnitude spectra of elementary sources in **V**
- ▶ **H** represents the *gains or activation functions* of the spectral profiles
  - i.e. in a given frame which elementary sources are active

## Toy example

- ▶ Two sine waves that come and go...





- ▶ NMF can be expressed as an optimization problem

$$\min_{\mathbf{W} \geq 0, \mathbf{H} \geq 0} \mathcal{D}(\mathbf{V}, \mathbf{W}\mathbf{H})$$

- ▶ The distance between the matrices is calculated element-wise

$$\mathcal{D}(\mathbf{V}, \mathbf{W}\mathbf{H}) = \sum_f \sum_n d(v_{fn}, w_{fk}h_{kn})$$

scalar distance

- ▶ Only  $\mathbf{V}$  is known.
  - NMF performs **unsupervised** learning

- ▶ Which distance to choose?
  - The most straightforward: Euclidean **distance**

$$d_{EUC}(x|y) = \frac{1}{2}(x - y)^2$$

- (Generalized) Kullback-Leibler **divergence**

$$d_{KL}(x|y) = x \log \frac{x}{y} - x + y$$

- Itakura-Saito **divergence**

$$d_{IS}(x|y) = \frac{x}{y} - \log \frac{x}{y} - 1$$

## Itakura-Saito distance

- ▶ Proposed by Itakura and Saito in 1968
- ▶ Measures the difference between two **spectra**
- ▶ Reflects **perceptual similarity** between two spectra
  - used as a speech enhancement performance metric
- ▶ It is **scale invariant**
  - low and high energy data are treated the same

$$d_{IS}(x|y) = d_{IS}(\alpha x|\alpha y)$$



- ▶ A statistical insight behind the choice of distance
  - Euclidean distance: ML estimation of  $W$ ,  $H$  in AGN

$$\mathbf{V} = \mathbf{WH} + \mathbf{E}$$

- KL divergence: ML estimation of  $W$ ,  $H$  in Poisson noise
- IS divergence: ML estimation of  $W$ ,  $H$  in Gamma multiplicative noise

$$\mathbf{V} = \mathbf{WH} \odot \mathbf{E}$$

- ▶ All the previous measures belong to the family of **beta divergences**:

$$d_{\beta}(x|y) = \begin{cases} \frac{x}{y} - \log \frac{x}{y} - 1 & \beta = 0 \\ x \log \frac{x}{y} + (y - x) & \beta = 1 \\ \frac{1}{\beta(\beta - 1)} (x^{\beta} + (\beta - 1)y^{\beta} - \beta xy^{\beta-1}) & \beta \in \mathbb{R} \setminus \{0, 1\} \end{cases}$$

- ▶ How can we **solve** the optimization problem?

$$\min_{\mathbf{W} \geq 0, \mathbf{H} \geq 0} \mathcal{D}(\mathbf{V}, \mathbf{WH})$$

- ▶ It cannot be solved for **both**  $\mathbf{W}$  and  $\mathbf{H}$ .
  - Keep  $\mathbf{W}$  **fixed** and optimize  $\mathbf{H}$  and vice versa
- ▶ Common minimization approach: **gradient descent**

$$\boldsymbol{\theta} \leftarrow \boldsymbol{\theta} - \eta \nabla \mathcal{J}(\boldsymbol{\theta})$$

- batch approach - all data are considered
- calculate the gradient of the distance w.r.t.  $\mathbf{W}$  and  $\mathbf{H}$
- rearrange terms in order to remove negative terms
- choose the appropriate step size



► Let's do this for **W**

- The derivative of the beta divergence

$$d_{\beta}(x|y) = y^{\beta-2}(y - x)$$

- The gradient with respect to **W**

$$\nabla_W D_{\beta}(\mathbf{V}|\mathbf{WH}) = \left( \mathbf{WH}^{\beta-2} \odot (\mathbf{WH} - \mathbf{V}) \right) \mathbf{H}^T$$

- Rearrange

$$\nabla_W D_{\beta}(\mathbf{V}|\mathbf{WH}) = \mathbf{W}^{\beta-1} \mathbf{H}^T \left( \mathbf{WH}^{\beta-2} \odot (\mathbf{WH} - \mathbf{V}) \right) \mathbf{H}^T$$

► Let's do this for **W**

- The gradient descent method

$$\mathbf{W}^{i+1} \leftarrow \mathbf{W}^i - \eta \left( \nabla_{\mathbf{W}}^+ D_{\beta}(\mathbf{V} | \mathbf{W} \mathbf{H}) - \nabla_{\mathbf{W}}^- D_{\beta}(\mathbf{V} | \mathbf{W} \mathbf{H}) \right)$$

- Choose the step size as

$$\eta = \frac{\mathbf{W}^i}{\nabla_{\mathbf{W}}^+ D_{\beta}(\mathbf{V} | \mathbf{W} \mathbf{H})}$$

- Substitute

$$\mathbf{W}^{i+1} \leftarrow \mathbf{W}^i \frac{\nabla_{\mathbf{W}}^- D_{\beta}(\mathbf{V} | \mathbf{W} \mathbf{H})}{\nabla_{\mathbf{W}}^+ D_{\beta}(\mathbf{V} | \mathbf{W} \mathbf{H})}$$

- NMF multiplicative update rules for beta divergence

$$\mathbf{W} \leftarrow \mathbf{W} \frac{\left( (\mathbf{W}\mathbf{H})^{\beta-2} \odot \mathbf{V} \right) \mathbf{H}^T}{(\mathbf{W}\mathbf{H})^{\beta-1} \mathbf{H}^T}$$

$$\mathbf{H} \leftarrow \mathbf{H} \frac{\mathbf{W}^T \left( (\mathbf{W}\mathbf{H})^{\beta-2} \odot \mathbf{V} \right)}{\mathbf{W}^T (\mathbf{W}\mathbf{H})^{\beta-1}}$$

- Why multiplicative?
  - non-negativity is preserved
  - easy to implement



- ▶ The MU updates are proved to **converge**
  - Lee & Seung's classic paper
- ▶ The NMF problem is **not convex for both  $W$  and  $H$** 
  - it is convex for  $W$  or  $H$
  - It is **not guaranteed** that a global minimum will be found

► The basic NMF algorithm

---

**Algorithm 1** NMF

**Require:**  $V \in \mathbb{R}_+, K$

Initialize  $W, H$

**for**  $i = 1$  **to** iterations **do**

    Update  $H$

    Update  $W$

    Normalize  $W, H$

**end for**

**return**  $W, H$

---

## How to start?

- ▶ Which are the initial values of  $\mathbf{W}$  and  $\mathbf{H}$ ?
  - the most common approach: **random non-negative values!**
    - the specific random value distribution may affect results
  - NMF is generally **sensitive** to initial values
  - more involved initialization strategies have been proposed
  - initialization **depends** on application
    - for audio applications, random is enough



## When to stop?

- ▶ NMF is calculated iteratively.
  - How **many** iterations do we need?
- ▶ Typically you check that the cost function decreases more than  $\varepsilon$  (tolerance) in each iteration
- ▶ You can also set a **maximum** number of iterations.
- ▶ For audio applications: trial and error
  - a smaller value of the cost function does not mean better perceptual quality

- ▶ How do we choose  $K$ ?
  - model order selection is a difficult problem
  - most approaches are trial and error
  - some approaches based on the eigenvalues of  $\mathbf{V}$
  - the number of components  $K$  is not necessarily equal to the number of sources
- ▶ Normalization
  - avoid getting stuck by small values due to multiplicative updates
  - normalize columns of  $\mathbf{W}$
  - the cost function changes!

- ▶ Adding constraints: the beauty of NMF
  - incorporate **prior knowledge** about the data
  - easy to extend the core method for specific applications
- ▶ We now want to solve

$$\min_{\mathbf{W} \geq 0, \mathbf{H} \geq 0} \mathcal{D}(\mathbf{W}, \mathbf{V}, \mathbf{H})$$

- ▶ where the cost function is

$$\mathcal{J}(\mathbf{W}, \mathbf{H}) = \mathcal{D}_{\beta}(\mathbf{V}, \mathbf{WH}) + c_1 \Phi(\mathbf{W}) + c_2 \Psi(\mathbf{H})$$



- The multiplicative updates have the same form:

$$\mathbf{W} \leftarrow \mathbf{W} \frac{\nabla_{\mathbf{W}}^{-} \mathcal{J}(\mathbf{W}, \mathbf{H})}{\nabla_{\mathbf{W}}^{+} \mathcal{J}(\mathbf{W}, \mathbf{H})}$$

- where each gradient is rearranged to form

$$\nabla_{\mathbf{W}}^{+} \mathcal{J}(\mathbf{W}, \mathbf{H}) = \nabla_{\mathbf{W}}^{+} \mathcal{D}_{\beta}(\mathbf{V}, \mathbf{WH}) + c_1 \nabla_{\mathbf{W}}^{+} \Phi(\mathbf{W}) + c_2 \cancel{\nabla_{\mathbf{W}}^{+} \Psi(\mathbf{H})}$$

$$\nabla_{\mathbf{W}}^{-} \mathcal{J}(\mathbf{W}, \mathbf{H}) = \nabla_{\mathbf{W}}^{-} \mathcal{D}_{\beta}(\mathbf{V}, \mathbf{WH}) + c_1 \nabla_{\mathbf{W}}^{-} \Phi(\mathbf{W}) + c_2 \cancel{\nabla_{\mathbf{W}}^{-} \Psi(\mathbf{H})}$$

- Of course the same holds for  $\mathbf{H}$

- ▶ The most common constraint: **sparsity!**
  - each data point is a combination of **some** sources
- ▶ How is sparsity measured?
  - There are several ways which can be summarized as

$$\Psi(\mathbf{H}) = \sum_{k,n} f(h_{kn})$$

- A very simple and efficient way is to choose  $f(x) = x$
- Since all elements are non-negative it is equivalent to the entrywise  $l_1$  norm

$$\Psi(\mathbf{H}) = \|\mathbf{H}\|_1$$

- So the cost function now is

$$\mathcal{J}(\mathbf{W}, \mathbf{H}) = \mathcal{D}_{\beta}(\mathbf{V}, \mathbf{WH}) + \lambda \|\mathbf{H}\|_1$$

- The gradients of the constraint are

$$\nabla_H^+ \Psi(\mathbf{H}) = 1 \quad \nabla_H^- \Psi(\mathbf{H}) = 0$$

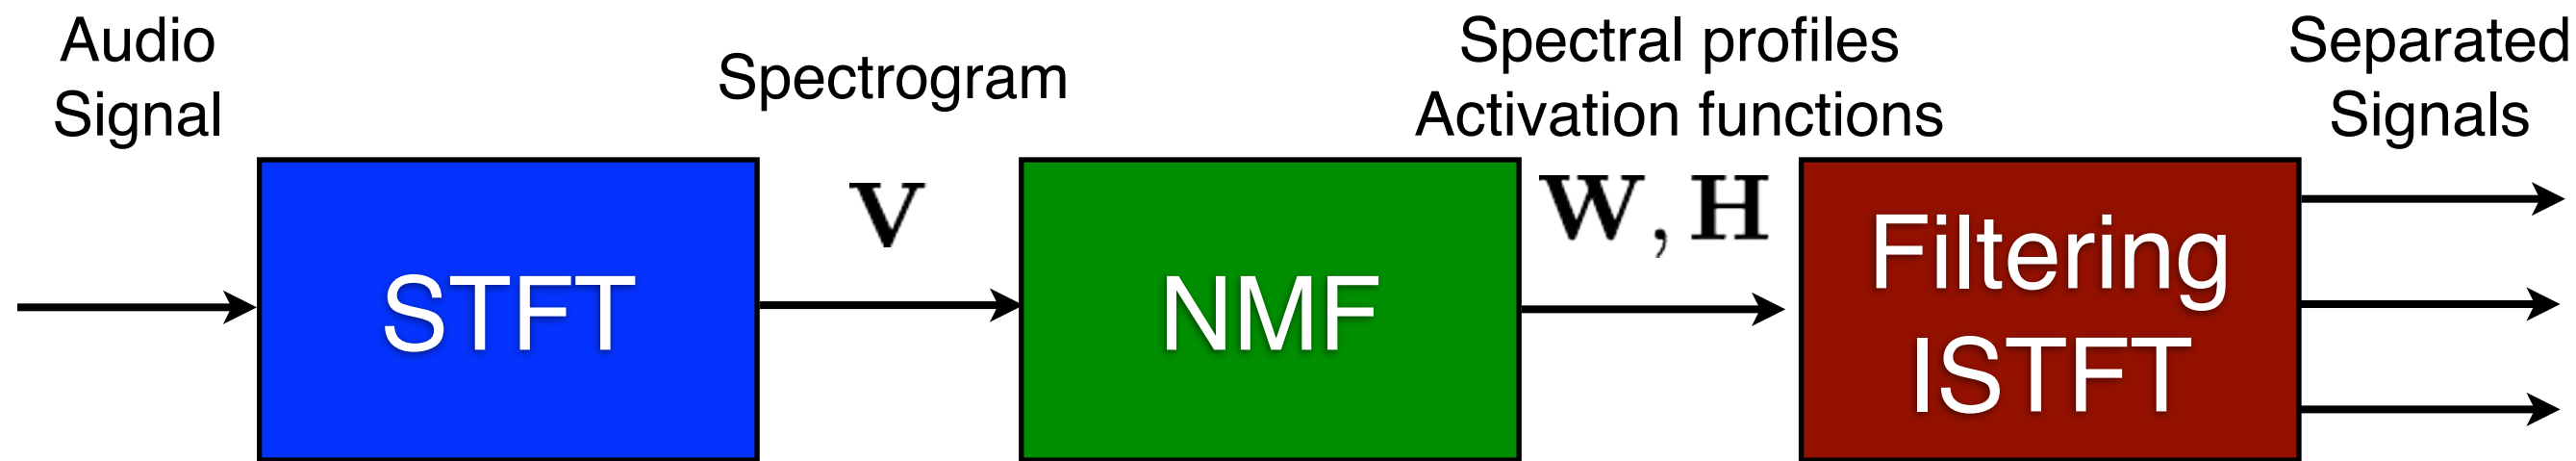
- And the new multiplicative update rules are

$$\mathbf{W} \leftarrow \mathbf{W} \frac{\left( (\mathbf{WH})^{\beta-2} \odot \mathbf{V} \right) \mathbf{H}^T}{(\mathbf{WH})^{\beta-1} \mathbf{H}^T} \quad \mathbf{H} \leftarrow \mathbf{H} \frac{\mathbf{W}^T \left( \mathbf{V} \odot (\mathbf{WH})^{\beta-2} \right)}{\mathbf{W}^T (\mathbf{WH})^{\beta-1}} + \lambda$$



# How to separate using NMF?

- Building blocks of an NMF based separation system



- ▶ STFT: from time to time-frequency domain
- ▶ Why? Audio signals are sparse in this domain
- ▶ Time-frequency resolution:
  - trade-off: choose good time OR frequency resolution
  - workaround: decrease hop size
    - this increases the number of data points tremendously
- ▶ Other things to consider:
  - analysis window type
  - spectrogram domain (magnitude or power)

## Separating with masks

- ▶ We know how to get **V**, **W** and **H**.
- ▶ How do we get the separated signals?
- ▶ Time-frequency masks

- Recall the Wiener filter

$$H(\omega) = \frac{P_{ss}(\omega)}{P_{ss}(\omega) + P_{nn}(\omega)}$$

- Pseudo-Wiener masks

$$\mathbf{M}_k = \frac{\mathbf{w}_k \mathbf{h}^k}{\mathbf{W} \mathbf{H}}$$

- Masks are real. Signal phase is left untouched.



## Separating with masks

- ▶ Masks are applied on the complex spectrogram **X**
- ▶ Each mask produces a new spectrogram
- ▶ We perform ISTFT to obtain the separated signal

# Implementation



- ▶ NMF can be trained!
  - Semi-supervised case
  - Supervised case
- ▶ Sort components: Make sense of the output data!
- ▶ Multichannel extensions: tensor factorizations
- ▶ Bayesian formulations
- ▶ Probabilistic Latent Component Analysis (PLCA)



# References

- ▶ D.D. Lee and H.B. Seung, “Learning the parts of objects by non-negative matrix factorization”, *Nature* **401**, 788-791
- ▶ D.D. Lee and H.B. Seung, “Algorithms for non-negative matrix factorization”, NIPS 2000
- ▶ N. Gillis, “*The Why and How of Nonnegative Matrix Factorization*”, In: “Regularization, Optimization, Kernels, and Support Vector Machines”, J.A.K. Suykens, M. Signoretto and A. Argyriou (eds), Chapman & Hall/CRC, Machine Learning and Pattern Recognition Series, pp. 257-291, 2014.
- ▶ Y.-X. Wang and Y.-J. Zhang, “Nonnegative Matrix Factorization: A Comprehensive Review,” *IEEE Trans. Knowl. Data Eng.*, vol. 25, no. 6, pp. 1336–1353.
- ▶ T. Virtanen, J. F. Gemmeke, B. Raj, and P. Smaragdis, “Compositional Models for Audio Processing: Uncovering the structure of sound mixtures,” *IEEE Signal Process. Mag.*, vol. 32, no. 2, pp. 125–144, Feb. 2015.
- ▶ T. Virtanen, “Monaural Sound Source Separation by Nonnegative Matrix Factorization With Temporal Continuity and Sparseness Criteria,” *IEEE Trans. Audio Speech Lang. Process.*, vol. 15, no. 3, pp. 1066–1074.
- ▶ A. Cichocki, R. Zdunek, A.H. Phan, “Non-negative matrix and tensor factorizations: applications to exploratory multi-way data analysis and blind source separation”

**Thank you!**

accu**u**sonus